

## Homework 6 Solutions

### Hansen 17.2

$E[e_{it}|X_{it}] = 0$  is not a strong enough condition for  $\hat{\beta}$  from a fixed effects regression to be unbiased. This condition says that  $e_{it}$  is (mean) independent of  $X_{it}$ , but it does not rule out that  $e_{it}$  could be related to, say,  $X_{it+1}$ . This is not an entirely strange case either, particularly if a good “shock” in the current period leads to the covariate changing in the next time period.

More specifically, recall that

$$\begin{aligned} E[\hat{\beta} - \beta | \mathbf{X}] &= \left( \sum_{i=1}^n \dot{\mathbf{X}}_i' \dot{\mathbf{X}}_i \right)^{-1} \sum_{i=1}^n \dot{\mathbf{X}}_i' E[\mathbf{e}_i | \mathbf{X}] \\ &= \left( \sum_{i=1}^n \dot{\mathbf{X}}_i' \dot{\mathbf{X}}_i \right)^{-1} \sum_{i=1}^n \dot{\mathbf{X}}_i' E[\mathbf{e}_i | \mathbf{X}_i] \end{aligned}$$

where  $\mathbf{X}$  is the  $nT \times k$  “data matrix” and the other notation is from class, and the second equality uses that the observations are independent of each other. Notice that,

$$E[\mathbf{e}_i | \mathbf{X}_i] = \begin{bmatrix} E[e_{i1} | X_{i1}, X_{i2}, \dots, X_{iT}] \\ E[e_{i2} | X_{i1}, X_{i2}, \dots, X_{iT}] \\ \vdots \\ E[e_{iT} | X_{i1}, X_{i2}, \dots, X_{iT}] \end{bmatrix}$$

The condition in the problem is not strong enough that this term is equal to 0. And, if it is some function of  $X$ , then  $\hat{\beta}$  would not, in general, be unbiased for  $\beta$ .

### Additional Question 2

Yes, they are correct. Notice that, because there are only two periods, this regression is equivalent to

$$\Delta Y_{i2} = \Delta \theta_2 + \alpha D_{i2} + \Delta e_{i2}$$

Moreover,

$$E[\Delta Y_2 | D_2 = 1] = \Delta \theta_2 + \alpha \tag{1}$$

$$E[\Delta Y_2 | D_2 = 0] = \Delta \theta_2 \tag{2}$$

which implies that

$$\begin{aligned}\alpha &= E[\Delta Y_2 | D_2 = 1] - E[\Delta Y_2 | D_2 = 0] \\ &= E[\Delta Y_2 | D_2 = 1] - E[\Delta Y_2(0) | D_2 = 0] \\ &= E[\Delta Y_2 | D_2 = 1] - E[\Delta Y_2(0) | D_2 = 1] \\ &= ATT\end{aligned}$$

where the first line subtracts Equation 2 from Equation 1, the second equality writes the second term in terms of potential outcomes, and the third equality holds by the parallel trends assumption.

## Additional Question 2

```
set.seed(1234)
library(randomForest)
data <- as.data.frame(haven::read_dta("jtrain_observational.dta"))

n <- nrow(data)
data$id <- 1:n
fold1 <- subset(data, (id%%2) == 0)
fold2 <- subset(data, (id%%2) == 1)

ml_att <- function(f1, f2) {

  # use f1 to estimate the first step models
  Dmod <- randomForest(as.factor(train) ~ age + educ + black + hisp +
                        married + re75 + unem75, data=f1)
  Ymod <- randomForest(re78 ~ age + educ + black + hisp +
                        married + re75 + unem75, data=f1)

  # get predictions with f2
  pscore <- predict(Dmod, newdata=f2, type="prob")[,2] # this gets p(d=1|x)
  out_reg <- predict(Ymod, newdata=f2)

  # compute att(k) with f2
  D <- f2$train
  p <- mean(D)
  Y <- f2$re78
  att1 <- mean(D/p*(Y-out_reg))
  att2 <- mean((1-D)/p * pscore/(1-pscore) * (Y-out_reg) )
}
```

```

    att <- att1-att2
    att
  }

# cross splitting
ml1 <- ml_att(fold1,fold2)
# reverse roles
ml2 <- ml_att(fold2,fold1)
# average
mean(c(ml1,ml2))

```

```
## [1] 1.580708
```

This is larger than what we estimated on the last homework (about 0.85) though I noticed that my estimates do move somewhat if I run the code multiple times.

### Additional Question 3

#### Part (a)

```

library(Matrix)
load("job_displacement_clean2.RData")
# drop already treated
data <- subset(data, first.displaced != 2001)
data <- droplevels(data)
Y <- data$learn
data$D <- 1*( (data$year >= data$first.displaced) & data$first.displaced != 0)
X <- model.matrix(~ as.factor(year) + D, data=data)
n <- length(unique(data$id))
tp <- length(unique(data$year))
iT <- matrix(rep(1,tp))
Dmat <- bdiag(replicate(n,iT,simplify=FALSE))
M <- Matrix::Diagonal(n*tp) - Dmat%*%solve(t(Dmat)%*%Dmat)%*%t(Dmat)
bet <- solve(t(X)%*%M%*%X) %*% t(X)%*%M%*%Y
bet

```

```

## 8 x 1 Matrix of class "dgeMatrix"
##           [,1]
## [1,]  5.29358972
## [2,]  0.09436799

```

```
## [3,] 0.18195412
## [4,] 0.26904810
## [5,] 0.28752717
## [6,] 0.35242722
## [7,] 0.38427488
## [8,] -0.23557976
```

Next, let's calculate the standard errors where we use that

$$\begin{aligned}\sqrt{n}(\hat{\beta} - \beta) &= E[\mathbf{X}'_i \mathbf{M}_i \mathbf{X}_i]^{-1} \frac{1}{\sqrt{n}} \sum_{i=1}^n \mathbf{X}'_i \mathbf{M}_i \mathbf{e}_i + o_p(1) \\ &\xrightarrow{d} N(0, \mathbf{V})\end{aligned}$$

where

$$\begin{aligned}\mathbf{V} &= E[\mathbf{X}'_i \mathbf{M}_i \mathbf{X}_i]^{-1} \mathbf{\Omega} E[\mathbf{X}'_i \mathbf{M}_i \mathbf{X}_i]^{-1} \\ &= E[\dot{\mathbf{X}}'_i \dot{\mathbf{X}}_i]^{-1} \mathbf{\Omega} E[\dot{\mathbf{X}}'_i \dot{\mathbf{X}}_i]^{-1}\end{aligned}$$

and

$$\begin{aligned}\mathbf{\Omega} &= E[\mathbf{X}'_i \mathbf{M}_i \mathbf{e}_i \mathbf{e}'_i \mathbf{M}_i \mathbf{X}_i] \\ &= E[\dot{\mathbf{X}}'_i \mathbf{e}_i \mathbf{e}'_i \dot{\mathbf{X}}_i]\end{aligned}$$

It's worth thinking about how to actually estimate these because  $\dot{\mathbf{X}}_i$  is a matrix rather than our usual case of it being a vector. First, notice that

$$\dot{\mathbf{X}}'_i \dot{\mathbf{X}}_i = \sum_{t=1}^T \dot{X}_{it} \dot{X}'_{it}$$

which is a  $k \times k$  matrix. Thus, the natural estimate of  $E[\dot{\mathbf{X}}'_i \dot{\mathbf{X}}_i]$  is

$$\frac{1}{n} \sum_{i=1}^n \sum_{t=1}^T \dot{X}_{it} \dot{X}'_{it} = \dot{\mathbf{X}}' \dot{\mathbf{X}} / n$$

which corresponds to exactly the same way that we estimated this type of term throughout the semester. Next,

$$\dot{\mathbf{X}}'_i \mathbf{e}_i = \sum_{t=1}^T X_{it} e_{it}$$

which is a  $k \times 1$  vector and so that

$$\dot{\mathbf{X}}_i' \mathbf{e}_i \mathbf{e}_i' \dot{\mathbf{X}}_i = \left( \sum_{t=1}^T \dot{X}_{it} e_{it} \right) \left( \sum_{t=1}^T \dot{X}_{it} e_{it} \right)'$$

and implies that we would estimate  $\boldsymbol{\Omega}$  by

$$\hat{\boldsymbol{\Omega}} = \frac{1}{n} \sum_{i=1}^n \left( \sum_{t=1}^T \dot{X}_{it} \hat{e}_{it} \right) \left( \sum_{t=1}^T \dot{X}_{it} \hat{e}_{it} \right)'$$

As far as I know, you can't play the same matrix algebra "trick" that we usually use here (in particular, recall that in the cross sectional case we could estimate  $\hat{\boldsymbol{\Omega}} = \frac{1}{n} \sum_{i=1}^n X_i X_i' \hat{e}_i^2$ , but that, for programming, it was often convenient to re-express this  $\hat{\boldsymbol{\Omega}} = \tilde{\mathbf{X}}' \tilde{\mathbf{X}}/n$  where a typical element of  $\tilde{\mathbf{X}}$  is given by  $X_i \hat{e}_i$ .) Anyway, the line below that uses the `rowsum` function is essentially just manually calculating  $\sum_{t=1}^T X_{it} \hat{e}_{it}$  and then using matrix algebra below it.

```
ehat <- as.numeric(Y - X%*%bet)
n <- length(unique(data$id))
dotX <- M%*%X
Q <- t(X) %*% dotX / n
dotXe <- rowsum(as.matrix(dotX*ehat), group=data$id)
Omeg <- t(dotXe)%*%dotXe/n
V <- solve(Q)%*%Omeg%*%solve(Q)
se <- sqrt(diag(V))/sqrt(n)
round(cbind.data.frame(bet=as.numeric(bet), se=se), 4)
```

```
##      bet      se
## 1  5.2936 0.1003
## 2  0.0944 0.0085
## 3  0.1820 0.0098
## 4  0.2690 0.0108
## 5  0.2875 0.0115
## 6  0.3524 0.0116
## 7  0.3843 0.0124
## 8 -0.2356 0.0257
```

Thus, we estimate that job displacement reduces earnings by about 23%. As a check, let's compare this to what we get from `fixest`.

```
library(fixest)
fe_reg <- feols(learn ~ as.factor(year) + D | id, data=data)
summary(fe_reg)
```

```
## OLS estimation, Dep. Var.: learn
## Observations: 18,928
## Fixed-effects: id: 2,704
## Standard-errors: Clustered (id)
##
##          Estimate Std. Error  t value  Pr(>|t|)
## as.factor(year)2003  0.094368   0.008533 11.05873 < 2.2e-16 ***
## as.factor(year)2005  0.181954   0.009756 18.65073 < 2.2e-16 ***
## as.factor(year)2007  0.269048   0.010758 25.00982 < 2.2e-16 ***
## as.factor(year)2009  0.287527   0.011507 24.98624 < 2.2e-16 ***
## as.factor(year)2011  0.352427   0.011598 30.38799 < 2.2e-16 ***
## as.factor(year)2013  0.384275   0.012396 30.99875 < 2.2e-16 ***
## D                    -0.235580   0.025677 -9.17486 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.352029      Adj. R2: 0.75681
##
##          Within R2: 0.107896
```

These appear to be the same (or the same up to possibly a degree of freedom adjustment).

## Part (b)

```
# list of time periods
# for simplicity I'm going to convert this to 1,2,3,4,5,6,7...
tlist <- (sort(unique(data$year)) - 1999)/2
# list of groups (excluding never-treated)
glist <- (sort(unique(data$first.displaced))[-1] - 1999)/2

# create new variables in updated time scale
data$G <- ifelse(data$first.displaced==0, 0, (data$first.displaced - 1999)/2)
data$tp <- (data$year-1999)/2

# write a function to compute att(g,t)
# I compute these as averages using weights, but it
# is fine to use subsets of data here too.
# @param w weights, used for bootstrap
# @param base_period, allows base period to optionally
# be fixed at one
compute.attgt <- function(data, w=rep(1,nrow(data)),
                           use_base_period_1=FALSE) {
```

```

# data frame to store results
results <- list()
counter <- 1

for (this_t in tlist[-1]) {
  for (this_g in glist) {
    base_period <- min(this_t-1, this_g-1)
    if (use_base_period_1) base_period <- 1
    G <- 1*(data$G==this_g)
    U <- 1*(data$G==0)
    pre <- 1*(data$tp == base_period)
    post <- 1*(data$tp == this_t)
    pg <- weighted.mean(data$G == this_g, w=w)
    pu <- weighted.mean(data$G == 0, w=w)
    ppre <- mean(pre)
    ppost <- mean(post)

    this_attgt <- weighted.mean(data$learn*G*post/pg/ppost, w=w) -
      weighted.mean(data$learn*G*pre/pg/ppre, w=w) -
      (weighted.mean(data$learn*U*post/pu/ppost, w=w) -
        weighted.mean(data$learn*U*pre/pu/ppre, w=w))
    results[[counter]] <- c(attgt=this_attgt, g=this_g, t=this_t)

    counter <- counter+1
  }
}

# convert to data frame
results <- as.data.frame(do.call("rbind", results))

results
}

results <- compute.attgt(data)
# print results
round(results[order(results$g, results$t),],4)

```

```

##      attgt g t
## 1 -0.2091 2 2

```

```
## 7 -0.1562 2 3
## 13 -0.1775 2 4
## 19 -0.2375 2 5
## 25 -0.2347 2 6
## 31 -0.2781 2 7
## 2 0.0117 3 2
## 8 -0.1138 3 3
## 14 -0.1124 3 4
## 20 -0.1677 3 5
## 26 -0.1698 3 6
## 32 -0.0313 3 7
## 3 0.0726 4 2
## 9 -0.0641 4 3
## 15 -0.1989 4 4
## 21 -0.3070 4 5
## 27 -0.2000 4 6
## 33 -0.2506 4 7
## 4 -0.0128 5 2
## 10 0.0036 5 3
## 16 -0.0603 5 4
## 22 -0.3184 5 5
## 28 -0.3115 5 6
## 34 -0.2210 5 7
## 5 0.0195 6 2
## 11 -0.1013 6 3
## 17 -0.0129 6 4
## 23 0.0565 6 5
## 29 -0.2505 6 6
## 35 -0.1633 6 7
## 6 0.1183 7 2
## 12 -0.0379 7 3
## 18 -0.0295 7 4
## 24 0.0205 7 5
## 30 -0.1143 7 6
## 36 -0.2056 7 7
```



## Part (c)

```
# function to compute att0
# ret_weights argument optionally returns the underlying
# weights rather than att0
compute.att0 <- function(attgt_results, w=rep(1,nrow(data)),
                          ret_weights=FALSE) {
  # overall attgt weights
  ever_treated <- which(data$G != 0)
  w <- w[ever_treated]
  pg <- sapply(glist, function(g) weighted.mean(data[ever_treated,]$G==g, w=w))
  maxT <- max(tlist)
  w0 <- function(g,t) {
    1*(t >= g)*pg[glist==g] / (maxT - g + 1)
  }
  # add weights to results
  w0gt <- sapply(1:nrow(attgt_results),
                 function(i) w0(attgt_results$g[i], attgt_results$t[i]))
  attgt_results$w0 <- w0gt
  # optionally return computed weights
  if(ret_weights) return(attgt_results)
  att0 <- sum(attgt_results$attgt*attgt_results$w0)
  att0
}

att0 <- compute.att0(results)

# bootstrap standard errors
B <- 100
id_list <- unique(data$id)
boot_att0 <- list()
for (b in 1:B) {
  # draw weights from multinomial distribution (this is exactly the same
  # as empirical bootstrap)
  boot_weights <- as.numeric(rmultinom(n=1, size=n, prob=rep(1/n,n)))
  this_boot_weights_id <- cbind.data.frame(id=id_list, boot_weights=boot_weights)
  boot_data <- merge(data, this_boot_weights_id, by="id")
  boot_attgt <- compute.attgt(data, w=boot_data$boot_weights)
  boot_att0[[b]] <- compute.att0(boot_attgt, w=boot_data$boot_weights)
```

```

}

boot_att0 <- do.call("rbind", boot_att0)
se <- sd(boot_att0)

round(cbind.data.frame(att0=att0, se=se), 4)

##      att0      se
## 1 -0.2111 0.0236

```

Thus, we are estimating a large, negative and statistically significant effect of job displacement.

## Part (d)

```

# twfe weights
Edt <- function(t) {
  mean( (t >= data$G) & (data$G !=0) )
}
mEdt <- mean(sapply(tlist, Edt))
pg2 <- sapply(glist, function(g) mean(data$G==g))
maxT <- max(tlist)
wTWFE_num <- function(g,t,post0=FALSE) {
  if ((t < g) & post0) return(0)
  hgt <- 1*(t>=g) - (maxT-g+1)/maxT - Edt(t) + mEdt
  hgt*pg2[glist==g]
}
# add weights to results
wTWFEgt_num <- sapply(1:nrow(results),
  function(i) wTWFE_num(results$g[i],
    results$t[i],
    post0=TRUE))
wTWFEgt_den <- sapply(1:nrow(results),
  function(i) wTWFE_num(results$g[i],
    results$t[i],
    post0=TRUE))
wTWFEgt <- wTWFEgt_num/sum(wTWFEgt_den)
results$wTWFE <- wTWFEgt

# get results from att0

```

```
results$watt0 <- compute.att0(results, ret_weights=TRUE)$w0
```

```
# print results
```

```
round(results[order(results$g, results$t),], 4)
```

```
##      attgt g t  wTWFE  watt0
## 1  -0.2091 2 2 0.0394 0.0378
## 7  -0.1562 2 3 0.0330 0.0378
## 13 -0.1775 2 4 0.0284 0.0378
## 19 -0.2375 2 5 0.0218 0.0378
## 25 -0.2347 2 6 0.0171 0.0378
## 31 -0.2781 2 7 0.0105 0.0378
## 2   0.0117 3 2 0.0000 0.0000
## 8  -0.1138 3 3 0.0465 0.0344
## 14 -0.1124 3 4 0.0430 0.0344
## 20 -0.1677 3 5 0.0380 0.0344
## 26 -0.1698 3 6 0.0344 0.0344
## 32 -0.0313 3 7 0.0294 0.0344
## 3   0.0726 4 2 0.0000 0.0000
## 9  -0.0641 4 3 0.0000 0.0000
## 15 -0.1989 4 4 0.0454 0.0303
## 21 -0.3070 4 5 0.0419 0.0303
## 27 -0.2000 4 6 0.0394 0.0303
## 33 -0.2506 4 7 0.0358 0.0303
## 4  -0.0128 5 2 0.0000 0.0000
## 10  0.0036 5 3 0.0000 0.0000
## 16 -0.0603 5 4 0.0000 0.0000
## 22 -0.3184 5 5 0.0836 0.0594
## 28 -0.3115 5 6 0.0799 0.0594
## 34 -0.2210 5 7 0.0748 0.0594
## 5   0.0195 6 2 0.0000 0.0000
## 11 -0.1013 6 3 0.0000 0.0000
## 17 -0.0129 6 4 0.0000 0.0000
## 23  0.0565 6 5 0.0000 0.0000
## 29 -0.2505 6 6 0.0718 0.0626
## 35 -0.1633 6 7 0.0682 0.0626
## 6   0.1183 7 2 0.0000 0.0000
## 12 -0.0379 7 3 0.0000 0.0000
## 18 -0.0295 7 4 0.0000 0.0000
```

```
## 24  0.0205 7 5 0.0000 0.0000
## 30 -0.1143 7 6 0.0000 0.0000
## 36 -0.2056 7 7 0.1178 0.1761
```

Notice that none of the TWFE weights are negative here though some of them do seem fairly different from the weights on  $ATT^O$ .

As a final side-comment, you might notice that

```
attTWFE <- sum(results$attgt*results$wTWFE)
attTWFE
```

```
## [1] -0.2139115
```

is not exactly equal to  $\alpha$  that we calculated earlier. There are two reasons for this. First, our expression for  $\alpha$  in terms of underlying  $ATT(g, t)$ 's relied on parallel trends actually holding; so if it does not, then we will not get exactly the same thing. Second, there is estimation error in  $ATT(g, t)$ ; that is, we have  $\widehat{ATT}(g, t)$  rather than  $ATT(g, t)$ , and the way to estimate this is not unique. Let me very quickly give show how you can recover  $\alpha$ . In the notes, the line right before relating  $\alpha$  to underlying  $ATT(g, t)$ 's was

$$\alpha = \sum_{t=2}^T \sum_{g \in \bar{G}} h(g, t) \left( E[Y_{it} - Y_{i1} | G = g] - E[Y_{it} - Y_{i1} | U = 0] \right) p_g \bigg/ \sum_{t=1}^T E[\ddot{D}_{it}^2]$$

We can use this to compute a “decomposition” of  $\alpha$  that will be equal to what we actually estimated.

```
# this computes "ATT(g,t)'s" using base period = 1
# everywhere which is analogous to above equation
results2 <- compute.attgt(data, use_base_period_1=TRUE)
# compute weights but allow for non-zero weights
# in pre-treatment periods
wTWFEgt_num2 <- sapply(1:nrow(results),
                      function(i) wTWFE_num(results$g[i],
                                              results$t[i],
                                              post0=FALSE))
wTWFEgt_den2 <- sapply(1:nrow(results),
                      function(i) wTWFE_num(results$g[i],
                                              results$t[i],
                                              post0=TRUE))
wTWFEgt2 <- wTWFEgt_num2/sum(wTWFEgt_den2)
# check if this delivers alpha
sum(results2$attgt*wTWFEgt2)
```

```
## [1] -0.2355798
```

which are now the same.

## Part (e)

```
# function to compute event studies
compute.es <- function(attgt_results, w=rep(1,nrow(data))) {
  # event study weights
  eseq <- sort(unique(attgt_results$t - attgt_results$g))
  es_res <- list()
  counter <- 1
  for (e in eseq) {
    this_keepers <- which( (attgt_results$t - attgt_results$g) == e)
    this_attgt <- attgt_results$attgt[this_keepers]
    pg <- sapply(attgt_results$g[this_keepers],
                 function(g) weighted.mean(data$G==g, w=w))
    pg <- pg / sum(pg)
    att_e <- sum(this_attgt*pg)
    es_res[[counter]] <- c(att_e=att_e, e=e)
    counter <- counter+1
  }
  # convert to data frame
  es_results <- as.data.frame(do.call("rbind", es_res))
  es_results
}

es_results <- compute.es(results)


# bootstrap event study
B <- 100
id_list <- unique(data$id)
boot_es <- list()
for (b in 1:B) {
  boot_weights <- as.numeric(rmultinom(n=1, size=n, prob=rep(1/n,n)))
  this_boot_weights_id <- cbind.data.frame(id=id_list, boot_weights=boot_weights)
  boot_data <- merge(data, this_boot_weights_id, by="id")
  boot_attgt <- compute.attgt(data, w=boot_data$boot_weights)
  boot_es[[b]] <- compute.es(boot_attgt, w=boot_data$boot_weights)$att_e
}
```

```

boot_es <- do.call("rbind", boot_es)
se <- apply(boot_es, 2, sd)
es_results$se <- se
es_results$ciL <- es_results$att_e - 1.96*es_results$se
es_results$ciU <- es_results$att_e + 1.96*es_results$se

library(ggplot2)
ggplot(data=es_results, mapping=aes(x=e,y=att_e)) +
  geom_line() +
  geom_point(size=1.5) +
  geom_line(aes(y=ciU), linetype="dotted") +
  geom_line(aes(y=ciL), linetype="dotted") +
  scale_x_continuous(breaks=seq(-5,5), labels=seq(-10,10,2)) +
  theme_bw()

```



The figure suggests that job displacement causes earnings to drop by, on average, about 20% and that this effect is quite persistent; it appears to be roughly the same 10 years following job displacement. If you look at the estimates in pre-treatment periods, with the exception of 10 years before job displacement, the estimates are fairly close to 0 (and not statistically different from 0)

suggesting that the parallel trends assumption is likely to be fairly reasonable in this application.